



**Co-funded by
the European Union**



Erasmus+
Enriching lives, opening minds.

Co-funded by the European Union. Views and opinions expressed are however those of the author(s) only and do not necessarily reflect those of the European Union or the European Education and Culture Executive Agency (EACEA). Neither the European Union nor EACEA can be held responsible for them.

2023-1-EL01-KA220-HED-000166184

Final Project, Option 1: Where does your CNN make its mistakes?

Software that reads handwritten numbers, such as postcodes on envelopes or account numbers on paper forms, must be judged not only on how often it is right but on where it goes wrong, because a confident misreading carries a real cost. In this project you build such a model and study the mistakes it makes.

In this project you build a convolutional neural network that recognises handwritten digits, then you study the mistakes it makes.

Central question: where does your CNN make its mistakes, and does one change to the model move them?

The data set is MNIST: 70,000 images of handwritten digits, each 28 by 28 pixels in greyscale, labelled 0 to 9. You download it yourself in the next cell.

This project draws on two earlier modules. In Module 5, we built and trained dense neural networks in Keras. In Module 6, we added convolutional layers, used Dropout and EarlyStopping to control overfitting, and read a confusion matrix to see which classes a model mixes up. Here you bring those skills together on a larger, more varied data set.

How this notebook is organised. The cells in Part A are already written for you. They handle the import, the download, and the reshaping that turns each flat row of 784 numbers back into a 28 by 28 image. From Part B onward, the cells are yours to complete. Each one carries comments describing what to do, along with a few hints, but not the finished code. Work through them in order.

Training on the full data set may take a few minutes per epoch on a laptop without a graphics card. This is expected.

Part A. Setup and data

We start with the imports. This cell loads everything the notebook needs, so we run it once at the top.

```
import numpy as np
import matplotlib.pyplot as plt

from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

import keras
from keras import layers
```

Loading the data

We download MNIST from OpenML. The call returns the images in **x** as 70,000 rows of 784 numbers, and the labels in **y** as text. We convert the labels to integers so the model can work with them. The download runs once and may take a moment.

```
X, y = fetch_openml("mnist_784", version=1, as_frame=False, return_X_y=True)
y = y.astype(int)
```

```
print("Images:", X.shape)
print("Labels:", y.shape, "| classes:", np.unique(y))
```

```
Images: (70000, 784)
Labels: (70000,) | classes: [0 1 2 3 4 5 6 7 8 9]
```

✓ Preparing the images

Each row of `X` is one image flattened into a line of 784 numbers. We reshape every row back into a 28 by 28 image with a single grayscale channel, and we scale the pixel values from the range 0 to 255 down to the range 0 to 1, which helps the network train. We then split the data into a training set and a test set, keeping the class balance with `stratify`, as we did in Module 3.

```
X = X.reshape(-1, 28, 28, 1).astype("float32") / 255.0

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)

print("Training images:", X_train.shape)
print("Test images:      ", X_test.shape)

# If training is slow on your machine, you may work with a smaller sample
# while you get everything running, for example:
# X_train, y_train = X_train[:10000], y_train[:10000]
```

```
Training images: (56000, 28, 28, 1)
Test images:      (14000, 28, 28, 1)
```

✓ A first look at the images

Before modelling, we look at a few training images with their labels, to confirm the data loaded correctly.

```
fig, axes = plt.subplots(1, 6, figsize=(9, 2))
for ax, image, label in zip(axes, X_train, y_train):
    ax.imshow(image.reshape(28, 28), cmap="gray")
    ax.set_title(str(label))
    ax.axis("off")
plt.tight_layout()
plt.show()
```



✓ Part B. Build and train your network

From here on, the cells are yours to complete.

First, build a convolutional neural network with the Keras Sequential API, the same style you used in Module 6. Recall the shape of one image: 28 by 28, with one channel.

```
# Build a small convolutional network with the Keras Sequential API.
#
# A sensible starting point has these layers, in order:
# - an Input layer with the image shape (28, 28, 1) [given below]
# - a Conv2D layer (try 32 filters, a 3 by 3 window, relu activation)
# - a MaxPooling2D layer (a 2 by 2 window)
# - a Dropout layer (0.3 is a reasonable value to start with)
# - a Flatten layer (moves from the 2D feature maps to a 1D vector)
# - a Dense hidden layer (try 64 units, relu)
# - a Dense output layer (one unit per class, softmax activation)

keras.utils.set_random_seed(42)

model = keras.Sequential([
    layers.Input(shape=(28, 28, 1)),

    layers.Conv2D(32, (3, 3), activation="relu"),
    layers.MaxPooling2D((2, 2)),
```

```

layers.Dropout(0.3),
layers.Flatten(),
layers.Dense(64, activation="relu"),
layers.Dense(10, activation="softmax"),
], name="cnn_project")

model.summary()

```

Model: "cnn_project"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 64)	346,176
dense_1 (Dense)	(None, 10)	650

Total params: 347,146 (1.32 MB)
Trainable params: 347,146 (1.32 MB)

✓ Compile the model

Compiling sets the optimiser, the loss function, and the metric to watch during training.

```

# Compile the model. You need to choose:
#   - an optimiser: "adam" is a solid default
#   - a loss function: the labels are integers 0 to 9, not one-hot vectors,
#     so use the crossentropy loss that works with integer labels
#     (the same one you used in Module 6)
#   - a metric to watch: "accuracy"

model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"],
)

```

✓ Train the model

Train the network, watching how it performs on data it has not learnt from. `EarlyStopping` halts training once the validation loss stops improving, so you do not waste time or overfit.

```

# Train the model. Use:
#   - validation_split=0.1 so 10% of the training data is held back for checking
#   - an EarlyStopping callback: monitor "val_loss", patience of about 3,
#     and restore_best_weights=True
#   - up to about 15 epochs, and batch_size=128

early_stopping = keras.callbacks.EarlyStopping(
    monitor="val_loss",
    patience=3, #wait 3 epochs without improvement before stopping
    restore_best_weights=True, #keep the values from the best epoch
)

history = model.fit(
    X_train, y_train,
    validation_split=0.1,
    batch_size=128,
    epochs=15,
    callbacks=[early_stopping],
    verbose=2,
)

```

```

Epoch 1/15
394/394 - 26s - 65ms/step - accuracy: 0.9092 - loss: 0.3206 - val_accuracy: 0.9636 - val_loss: 0.1290
Epoch 2/15
394/394 - 43s - 110ms/step - accuracy: 0.9681 - loss: 0.1091 - val_accuracy: 0.9727 - val_loss: 0.0860
Epoch 3/15
394/394 - 39s - 98ms/step - accuracy: 0.9763 - loss: 0.0792 - val_accuracy: 0.9796 - val_loss: 0.0679
Epoch 4/15
394/394 - 24s - 61ms/step - accuracy: 0.9811 - loss: 0.0626 - val_accuracy: 0.9807 - val_loss: 0.0614

```

```

Epoch 5/15
394/394 - 24s - 62ms/step - accuracy: 0.9835 - loss: 0.0538 - val_accuracy: 0.9823 - val_loss: 0.0556
Epoch 6/15
394/394 - 41s - 105ms/step - accuracy: 0.9851 - loss: 0.0452 - val_accuracy: 0.9834 - val_loss: 0.0524
Epoch 7/15
394/394 - 23s - 59ms/step - accuracy: 0.9876 - loss: 0.0388 - val_accuracy: 0.9868 - val_loss: 0.0495
Epoch 8/15
394/394 - 43s - 109ms/step - accuracy: 0.9886 - loss: 0.0351 - val_accuracy: 0.9859 - val_loss: 0.0447
Epoch 9/15
394/394 - 25s - 62ms/step - accuracy: 0.9908 - loss: 0.0303 - val_accuracy: 0.9864 - val_loss: 0.0461
Epoch 10/15
394/394 - 23s - 59ms/step - accuracy: 0.9912 - loss: 0.0267 - val_accuracy: 0.9873 - val_loss: 0.0439
Epoch 11/15
394/394 - 43s - 108ms/step - accuracy: 0.9922 - loss: 0.0243 - val_accuracy: 0.9879 - val_loss: 0.0433
Epoch 12/15
394/394 - 25s - 63ms/step - accuracy: 0.9936 - loss: 0.0200 - val_accuracy: 0.9882 - val_loss: 0.0448
Epoch 13/15
394/394 - 41s - 104ms/step - accuracy: 0.9940 - loss: 0.0191 - val_accuracy: 0.9866 - val_loss: 0.0497
Epoch 14/15
394/394 - 41s - 104ms/step - accuracy: 0.9943 - loss: 0.0174 - val_accuracy: 0.9871 - val_loss: 0.0465

```

✓ Learning curves

Plot the training and validation loss across the epochs. A validation curve that flattens or turns upward while the training curve keeps falling is a sign of overfitting.

```

# Plot the training and validation loss from history.history.
# The keys you need are "loss" and "val_loss".


fig, (ax_loss, ax_acc) = plt.subplots(1, 2, figsize=(11,4))

ax_loss.plot(history.history["loss"], label="Training")
ax_loss.plot(history.history["val_loss"], label="Validation")
ax_loss.set_xlabel("Epoch"); ax_loss.set_ylabel("Loss")
ax_loss.set_title("Loss"); ax_loss.legend()

ax_acc.plot(history.history["accuracy"], label="Training")
ax_acc.plot(history.history["val_accuracy"], label="Validation")
ax_acc.set_xlabel("Epoch"); ax_acc.set_ylabel("Accuracy")
ax_acc.set_title("Accuracy"); ax_acc.legend()

plt.tight_layout()
plt.show()

```



✓ Part C. Read the mistakes

Now measure how well the model does on the test set, the images it has never seen. This gives a reliable estimate of how it will perform on new images.

```

# Evaluate the trained model on the test set, and produce predictions.
# - model.evaluate(X_test, y_test) returns the loss and the accuracy; store
#   them as test_loss and test_accuracy so you can compare the models later
# - model.predict(X_test) returns a probability for each of the 10 classes;
#   the predicted digit is the class with the highest probability,
#   which you can get with .argmax(axis=1)

```

```
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"Test accuracy: {test_accuracy: .3f}")
print(f"Test loss: {test_loss: .3f}")

y_pred_probs = model.predict(X_test, verbose=0) #a score for each digit
y_pred = y_pred_probs.argmax(axis=1) #the digit with the highest score
```

```
Test accuracy: 0.986
Test loss: 0.049
```


✓ The confusion matrix

The confusion matrix is where you answer the first half of the central question: where the mistakes are. Read it by rows. Each row is a true digit, each column is the digit the model predicted. The diagonal holds the correct predictions; every off-diagonal cell is a mistake, and its position tells you which digit was mistaken for which.

```
# Build and display the confusion matrix comparing y_test with y_pred.
# Use confusion_matrix(...) to compute it and ConfusionMatrixDisplay(...) to draw it.

cm = confusion_matrix(y_test, y_pred)

disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=range(10))
disp.plot(cmap="Blues", colorbar=False)
plt.title("Confusion Matrix on Model 1")
plt.show()
```



What did you find?

Look at the off-diagonal cells with the largest counts. Which pairs of digits does your model confuse most often? Some pairs are visually similar and are commonly mixed up.

Write two or three sentences here naming the two most common confusion pairs, recording the count for each one, and suggesting why those particular digits might be hard to tell apart. For example, if row 4 and column 9 contains 18, the model predicted 9 for 18 images that were actually 4.

The two most common confusion pairs are 4 misclassified as 9 (24 cases) and 2 misclassified as 8 (10 cases). The 4/9 confusion is understandable because both digits share a similar closed loop at the top, and handwritten versions can look nearly identical depending on writing style. The 2/8 confusion likely arises because of the upper and bottom curves existing on both numbers.

✓ Part D. One change

Now answer the second half of the central question: does one change to the model move the mistakes?

Choose **one** change from the two below. Do not make both at once, or you will not know which one caused any difference you see.

- **Option A: a deeper network.** Before the Flatten layer, add `Conv2D(64, (3, 3), activation="relu")` followed by `MaxPooling2D((2, 2))`. This lets the network learn more detailed patterns.

- **Option B: stronger regularisation.** Raise the Dropout rate from 0.3 to 0.5. This makes the network rely less on any single feature.

Before you run anything, write down your prediction: will your chosen change raise the test accuracy, and will it reduce the confusion between the pairs you found in Part C? Note it here.

Neural-network training contains some randomness. A very small difference in accuracy does not necessarily show that one model is genuinely better. Look at the size of the difference and at the specific mistakes, not only at which accuracy is larger.

First of all, I chose Option A with the addition of a second Conv2D layer before the Flatten one. I predict this addition will slightly increase the test accuracy compared to the previous one, as the deeper network can learn more about the details of the images. I also expect the 4/9 and 2/8 confusions to decrease as the second convolutional layer may capture more vital differences between these two pairs.

```
# Build a second model that applies your ONE chosen change.
# Start from your Part B model and change only what Option A or Option B asks for.
#
# Use the same seed as before to make the comparison more reproducible.
```

```
keras.utils.set_random_seed(42)
```

```
model2 = keras.Sequential([
    layers.Input(shape=(28, 28, 1)),

    layers.Conv2D(32, (3, 3), activation="relu"),
    layers.MaxPooling2D((2, 2)),

    #second convolution and pooling
    layers.Conv2D(64, (3, 3), activation="relu"),
    layers.MaxPooling2D((2, 2)),

    layers.Dropout(0.3),
    layers.Flatten(),
    layers.Dense(64, activation="relu"),
    layers.Dense(10, activation="softmax"),
], name="cnn_project_2")
```

```
model2.summary()
```

Model: "cnn_project_2"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
flatten_1 (Flatten)	(None, 1600)	0
dense_2 (Dense)	(None, 64)	102,464
dense_3 (Dense)	(None, 10)	650

Total params: 121,930 (476.29 KB)

Trainable params: 121,930 (476.29 KB)

```
# Compile and train model2 exactly as you trained the first model,
# so the comparison is fair (same optimiser, loss, metric, callback, and settings).
```

```
model2.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"],
)

early_stopping2 = keras.callbacks.EarlyStopping(
    monitor="val_loss",
    patience=3, #wait 3 epochs without improvement before stopping
    restore_best_weights=True, #keep the values from the best epoch
)

history2 = model2.fit(
    X_train, y_train,
    validation_split=0.1 ,
```

```

batch_size=128,
epochs=15,
callbacks=[early_stopping2],
verbose=2,
)

```

```

Epoch 1/15
394/394 - 47s - 119ms/step - accuracy: 0.9163 - loss: 0.2832 - val_accuracy: 0.9773 - val_loss: 0.0787
Epoch 2/15
394/394 - 42s - 106ms/step - accuracy: 0.9759 - loss: 0.0785 - val_accuracy: 0.9805 - val_loss: 0.0600
Epoch 3/15
394/394 - 81s - 207ms/step - accuracy: 0.9826 - loss: 0.0554 - val_accuracy: 0.9841 - val_loss: 0.0478
Epoch 4/15
394/394 - 41s - 105ms/step - accuracy: 0.9864 - loss: 0.0446 - val_accuracy: 0.9875 - val_loss: 0.0399
Epoch 5/15
394/394 - 41s - 104ms/step - accuracy: 0.9878 - loss: 0.0377 - val_accuracy: 0.9879 - val_loss: 0.0420
Epoch 6/15
394/394 - 41s - 104ms/step - accuracy: 0.9894 - loss: 0.0330 - val_accuracy: 0.9882 - val_loss: 0.0374
Epoch 7/15
394/394 - 41s - 105ms/step - accuracy: 0.9910 - loss: 0.0288 - val_accuracy: 0.9880 - val_loss: 0.0388
Epoch 8/15
394/394 - 82s - 208ms/step - accuracy: 0.9916 - loss: 0.0255 - val_accuracy: 0.9893 - val_loss: 0.0398
Epoch 9/15
394/394 - 41s - 104ms/step - accuracy: 0.9929 - loss: 0.0222 - val_accuracy: 0.9889 - val_loss: 0.0407

```

```

# Compare the two models.
# - store the results as test_loss2 and test_accuracy2
# - print test_accuracy2 next to the first model's test_accuracy
# - build y_pred2 and draw the new confusion matrix
# - compare the counts for the two confusion pairs you identified in Part C

test_loss2, test_accuracy2 = model2.evaluate(X_test, y_test, verbose=0)
print(f"Model 1 test accuracy: {test_accuracy: .3f}")
print(f"Model 2 Test accuracy: {test_accuracy2: .3f}")

y_pred_probs = model2.predict(X_test, verbose=0) #a score for each digit
y_pred2 = y_pred_probs.argmax(axis=1) #the digit with the highest score

cm2 = confusion_matrix(y_test, y_pred2)


disp2 = ConfusionMatrixDisplay(confusion_matrix=cm2, display_labels=range(10))
disp2.plot(cmap="Blues", colorbar=False)
plt.title("Confusion Matrix on Model 2")
plt.show()

```

```

Model 1 test accuracy: 0.986
Model 2 Test accuracy: 0.989

```



Ethical considerations

A digit reader does not perform equally well for everyone. Handwriting varies by age, country, and schooling, and a model trained mostly on one style may read another less reliably. **Write two or three sentences here** on the risks of putting this model into a real reading pipeline. Consider whose handwriting might be under-represented in the training data, and what the consequence of a confident but wrong reading would be in the setting above.

As we know the MNIST dataset was mainly collected from American writers and as we live in a multicultural world, the handwriting styles vary across countries, age groups and levels of education. In a real postal or banking system, even a small number of wrong predictions could cause chaos in the whole system. A safer approach would be a two factor validation where a human checks the flagged, low confidence predictions.

Wrapping up

Bring your findings together. **Write a short paragraph** that answers the central question directly:

- Where did your first model make most of its mistakes?
- Did your prediction in Part D turn out to be correct?
- For each of the two confusion pairs you identified, what was the count before and after the change?
- Did your one change move the mistakes, and is the difference large enough to be meaningful?

What to hand in: this completed notebook, with both confusion matrices visible in the output, and your written answers in the three markdown cells above (Part C, Part D, and this one).

As we mentioned before, the first model confused most often 4 with 9 (24 cases) and 2 with 8 (10 cases), because these digit pairs look similar in handwriting. My prediction was correct, adding a second Conv2D layer increased accuracy from 98.6% to 98.9%, and both confusion pairs improved: 4/9 dropped from 24 to 13 cases and 2/8 dropped from 10 to 7 cases. The accuracy difference is small, but the fact that both confusion pairs got better suggests the deeper network learned to tell these digits apart more reliably.